

Messaging Standards and Systems

AMQP & RabbitMQ

Gavin M. Roy
VP of Architecture
AWeber Communications
Twitter: @Crad

POSSCON

April 14, 2015

About Me

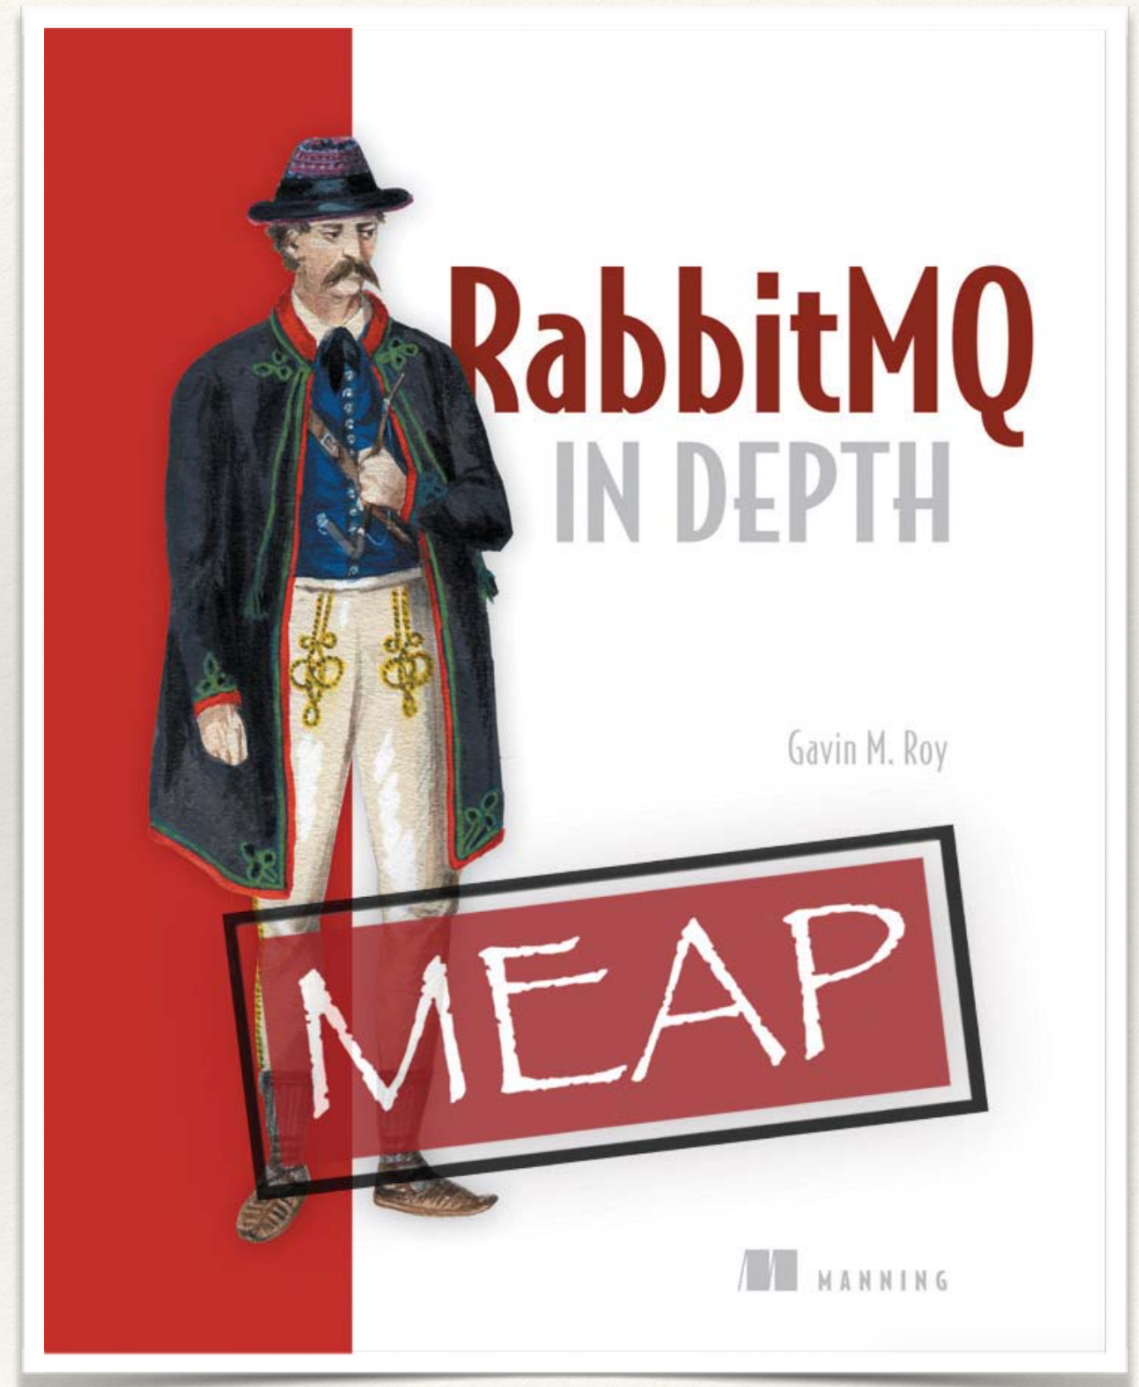
VP of Architecture
AWeber Communications

Blame me for pika, rabbitpy,
pamqp, and a handful of
RabbitMQ plugins

Blog: <https://gavinroy.com>

Github: <https://github.com/gmr>

Book: <http://manning.com/roy>

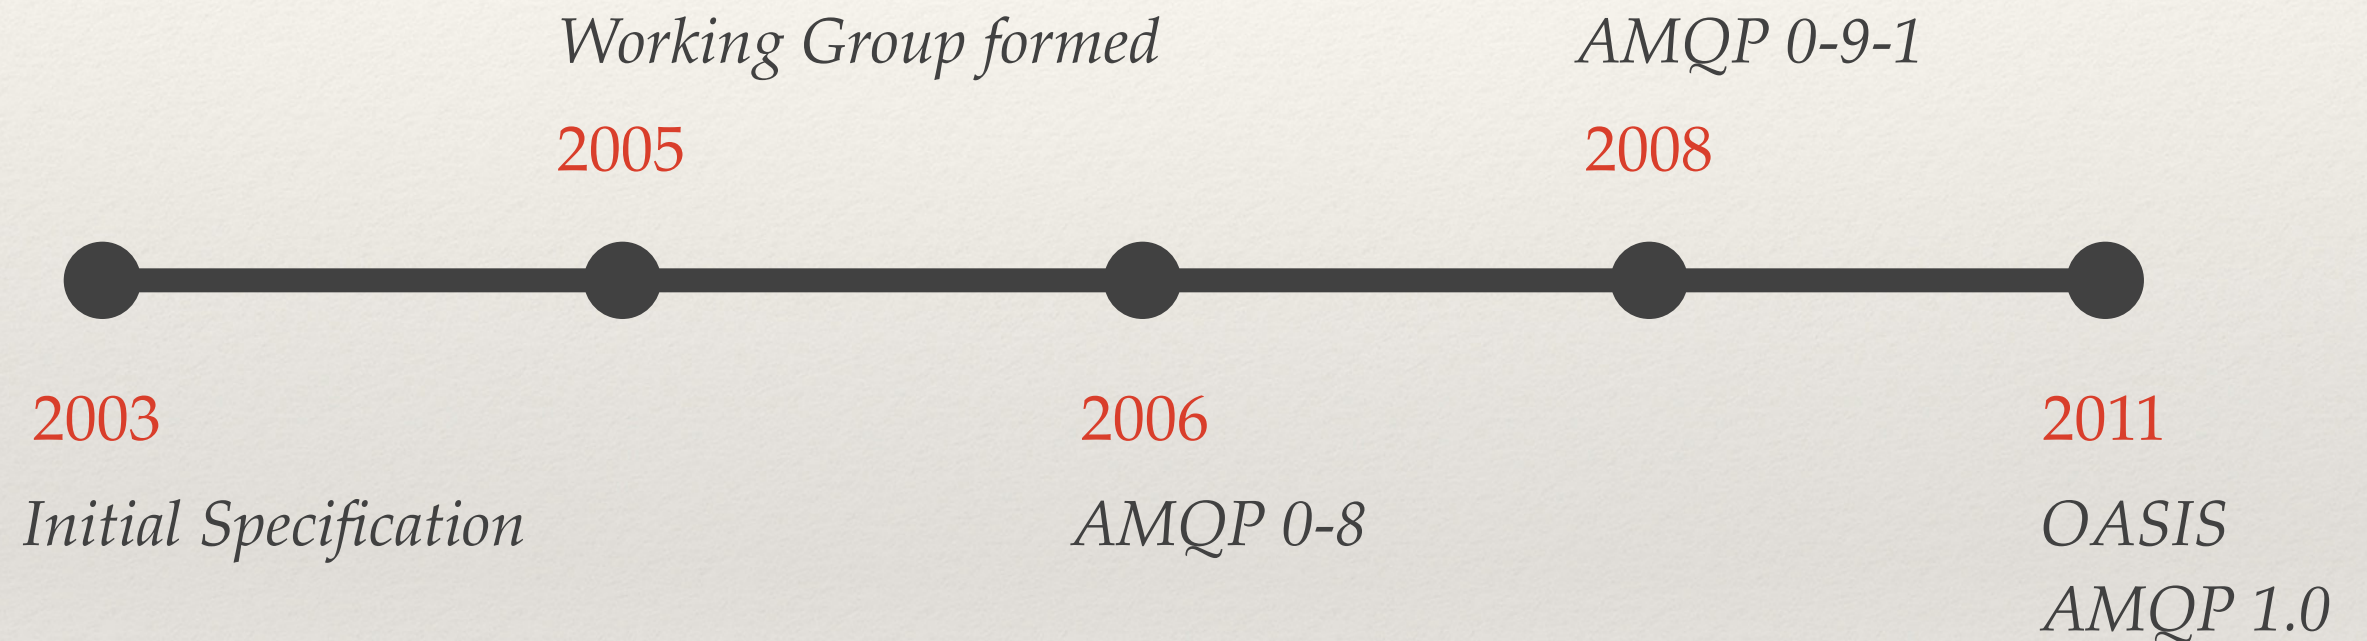


Advanced Message Queueing Protocol



- ❖ Open Standard
- ❖ Platform and Vendor Neutral
- ❖ Multiple Versions

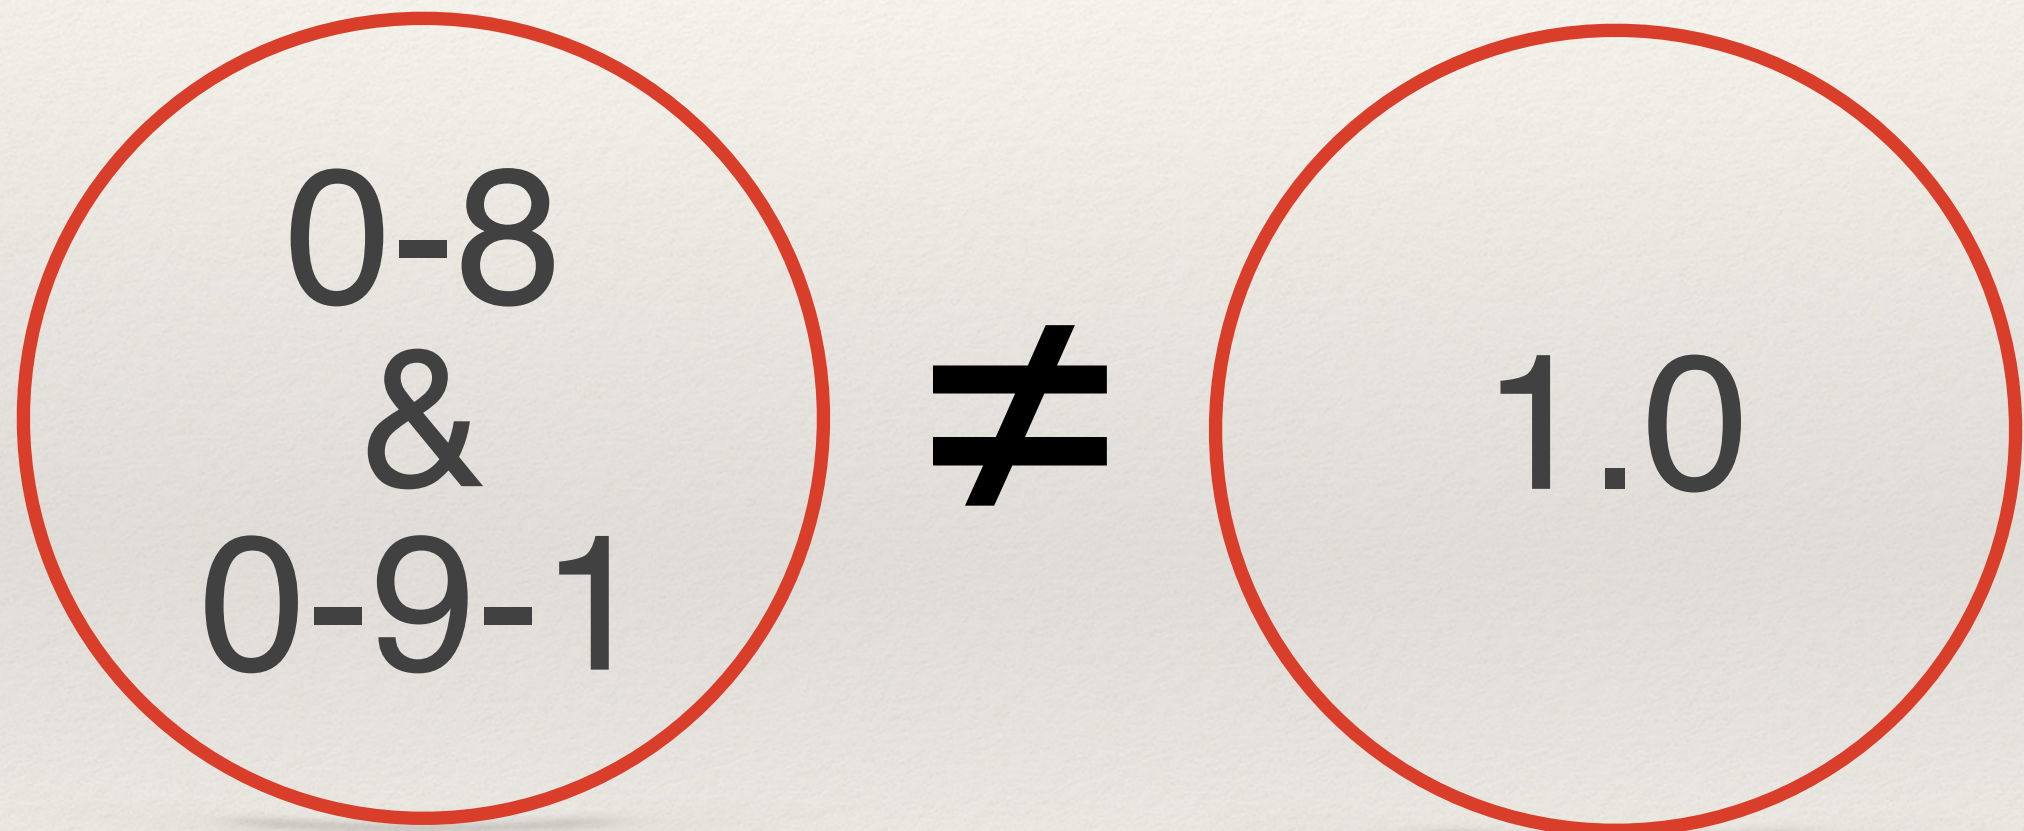
AMQP Timeline



AMQP Working Group

- ❖ JPMorgan Chase
- ❖ Deutsche Börse Systems
- ❖ Novell
- ❖ Cisco Systems
- ❖ Goldman Sachs
- ❖ Solace Systems
- ❖ IONA Technologies
- ❖ HCL Technologies
- ❖ Tervala, Inc.
- ❖ iMatrix
- ❖ Progress Software
- ❖ VMWare
- ❖ RedHat
- ❖ IIT Software
- ❖ WSO₂
- ❖ TWIST
- ❖ INETCO Systems Ltd.
- ❖ Bank of America
- ❖ Informatica Corporation
- ❖ Barclays
- ❖ Microsoft Corporation
- ❖ Credit Suisse
- ❖ my-Channels

Competing AMQP Standards



AMQP 0-9-1

- ❖ Currently has wider support than AMQP 1.0
- ❖ Multiple Broker Implementations:
RabbitMQ, Apache Qpid, and SwiftMQ to name a few
- ❖ Specifies a Model and Protocol

Native AMQP 0-8/0-9-1 Clients

C

Go

OCaml

Clojure

Groovy

Perl

Cobol

Haskell

PHP

Common Lisp

Java

Python

Delphi

JavaScript

Ruby

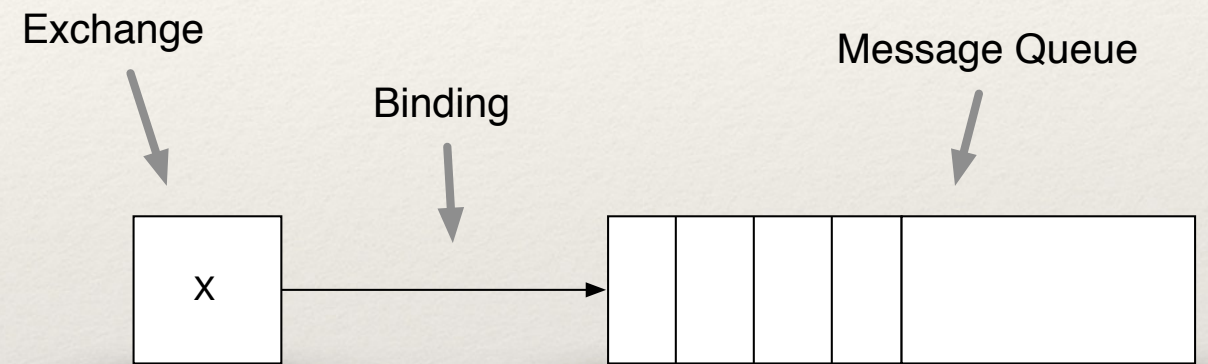
Erlang

.NET

Scala

Advanced Message Queueing Model

- ❖ **Exchange**
Receives and route messages
- ❖ **Message Queue**
Stores messages until they can be consumed
- ❖ **Binding**
Defines the relationship between an Exchange and Queue and provides routing criteria



Routing Keys

- ❖ Provided when publishing a message
- ❖ Compared against binding keys by exchanges
- ❖ Ideally provide context to the message:
 - ❖ *Connote the type of the message*
 - ❖ *Categorize the content in the message*
 - ❖ *Specify the type of consumer that should receive it*

Advanced Message Queueing Protocol

- ❖ Compact, binary frame format wire protocol
- ❖ Bi-directional RPC
- ❖ Commands consist of Classes and Methods:

- ❖ *Example Request:*

- `Queue.Declare(name="foo")`

- ❖ *Response:*

- `Queue.DeclareOk(messages=0, consumers=0)`

Common AMQP Terms

- ❖ Broker

A server that implements AMQP

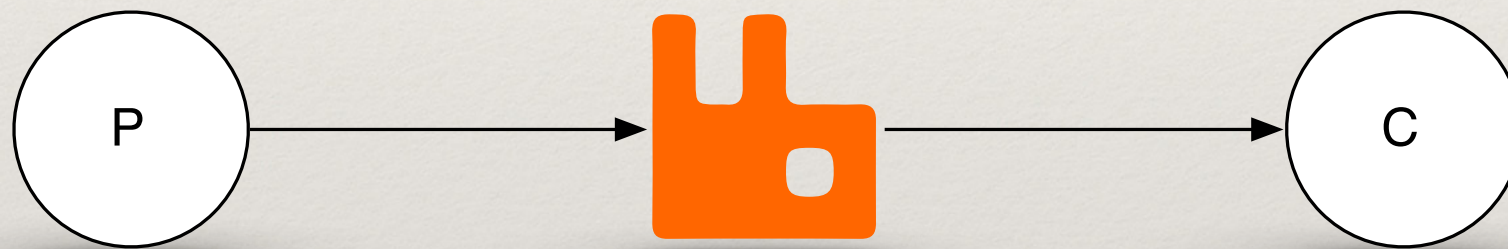
- ❖ Producer or Publisher

A client application that sends messages to a broker

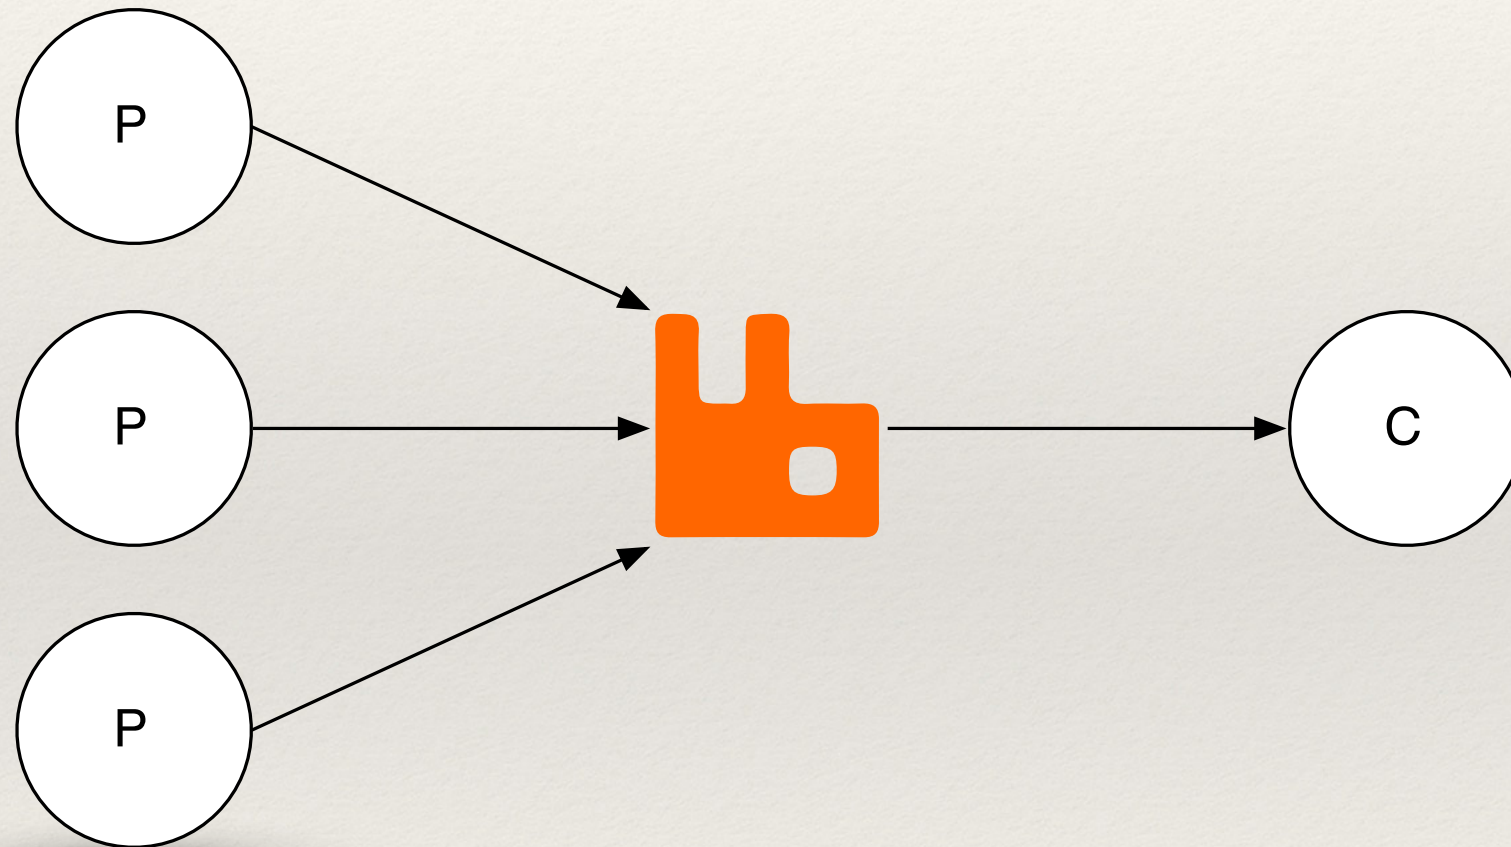
- ❖ Consumer

A client application that reads messages from a queue

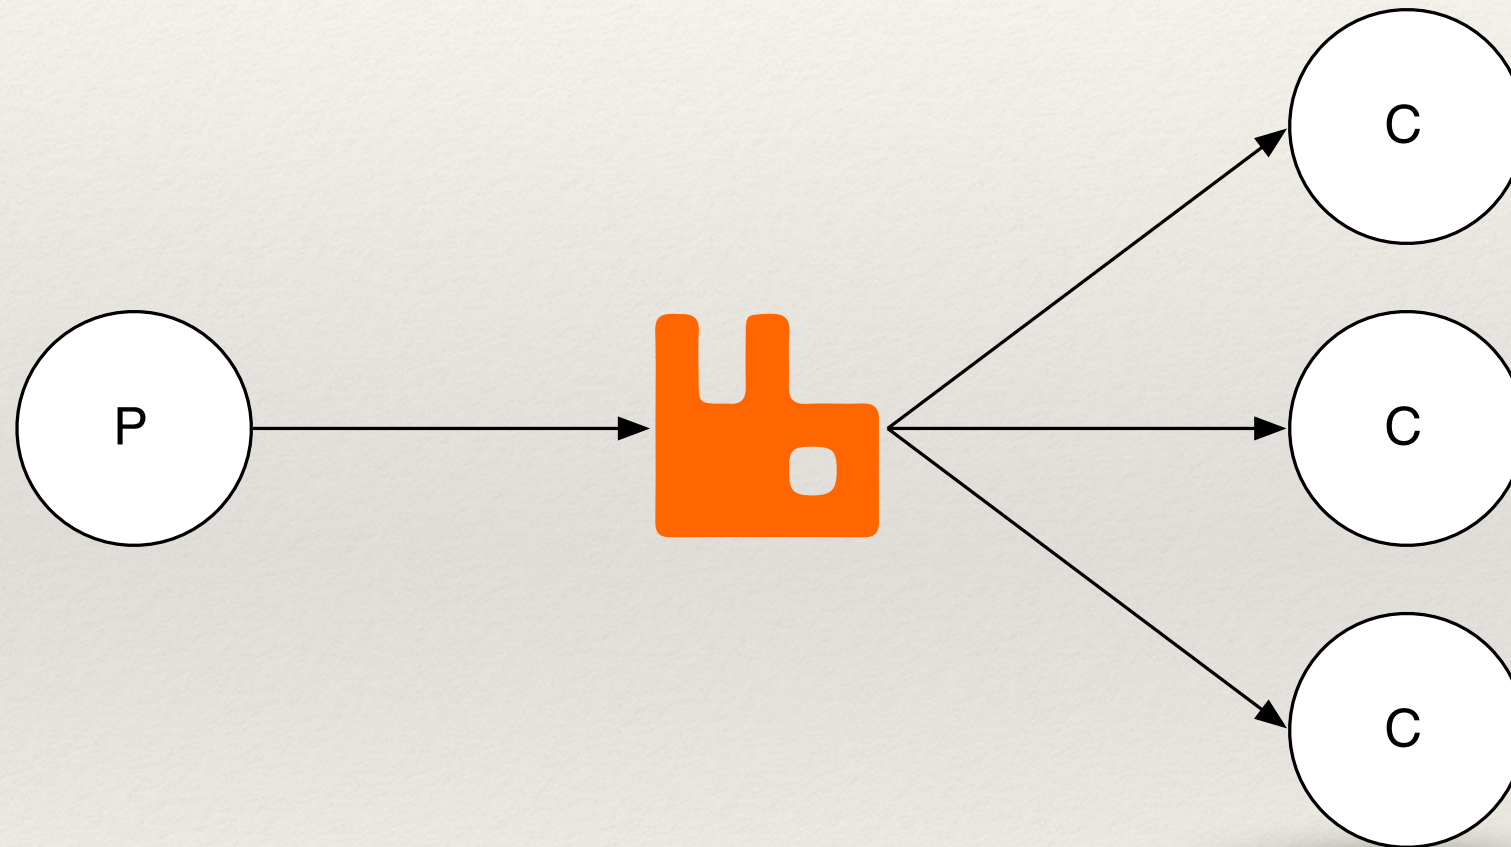
Publishers and Consumers



Multiple Publishers



Multiple Consumers



AMQP Messages

- ❖ Comprised of 3 or more frames:
 - ❖ **Method Frame**
Basic.Publish, Basic.Deliver, etc
 - ❖ **Content Header** with body size & message properties
timestamp, message-id, app-id, etc
 - ❖ *n* **Body** Frames with the opaque message payload

AMQP 0-9-1 Issues & Gotchas

- ❖ Ambiguous
- ❖ Authentication
- ❖ Asynchronous
- ❖ Connection Negotiation
- ❖ Exceptions



RabbitMQ



- ❖ Open Source (MPL)
- ❖ Written in Erlang / OTP
- ❖ Developed / Maintained by Pivotal
- ❖ Multi-Protocol
AMQP 0-9-1 & 1.0, MQTT, STOMP, XMPP, HTTP, Web-STOMP & More
- ❖ Roots in AMQP 0-8 / 0-9-1

Who Uses It?

Agora Games

NASA

Chef

New York Times

Google AdMob

National Science Foundation

Instagram

Openstack

MeetMe

Rapportive

Mercado Libre

Reddit

Mozilla

Soundcloud

(and many more)

Why Use RabbitMQ?

- ❖ Create loosely-coupled applications
- ❖ Communicate across applications or platforms
- ❖ Tap into pre-existing message flows for new purposes
- ❖ Scale-out clustering for growth, throughput, and HA
- ❖ Federation for WAN latencies and network partitions
- ❖ Extensible plugin-in architecture

RabbitMQ Extensions to AMQP

Authentication Failure

Connection blocking

Exchange to Exchange Bindings

Message CC & BCC Routing

Delivery Confirmations

Queue Length Limits

Basic.Nack

Per Queue Message TTL

Consumer cancellations

Per Message TTL

Consumer priorities

Queue TTL

Dead Letter Exchanges

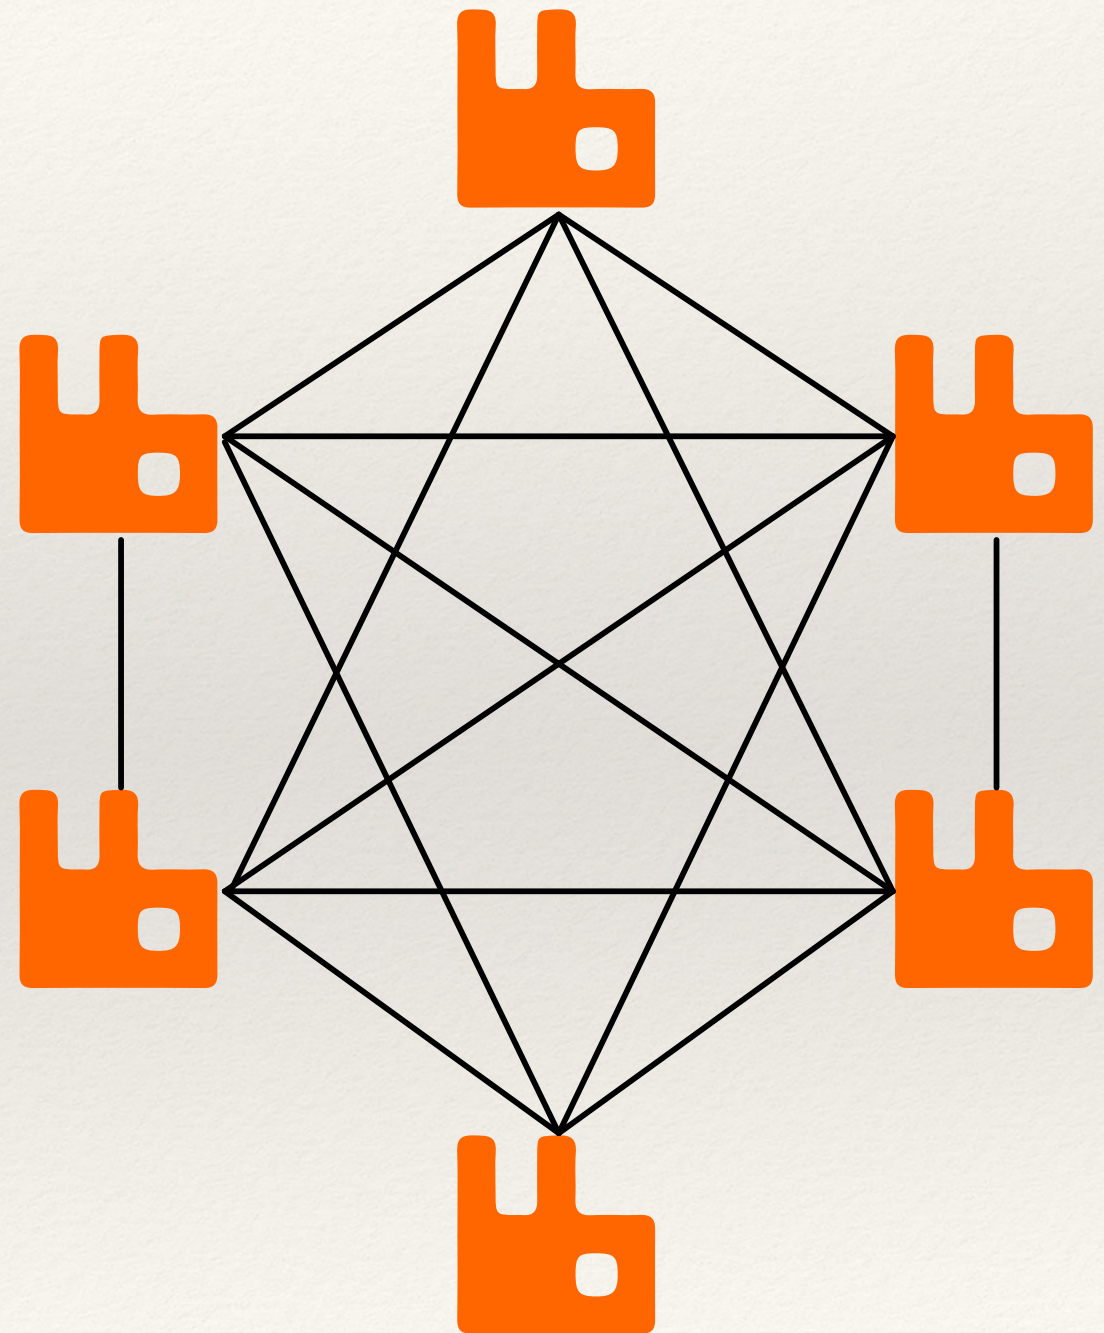
Message User ID Validation

Alternate Exchanges

Auto-delete exchanges


RabbitMQ Clustering

- ❖ LAN Only
- ❖ Adds highly-available queues
- ❖ Is cohesive, publish and consume from any node
- ❖ Leverages native Erlang clustering and communication
- ❖ Has multiple strategies for dealing with network partitions
- ❖ Manually configured via configuration or command line*



RabbitMQ Federation

- ❖ Provides loosely coupled inter-node/cluster RabbitMQ communication
- ❖ Connect multiple RabbitMQ clusters across data centers or the Internet
- ❖ Can be applied to both exchanges and queues



RabbitMQ Plugins

amqp1_0

Adds AMQP 1.0 support to RabbitMQ

auth-backend-amqp

AMQP RPC for Authentication

auth-backend-ldap

Authenticate via an external LDAP server

autocluster-consul

Automatically create clusters using Consul

msg_store_eleveldb_index

LevelDB storage for queue messages

presence-exchange

Publishes messages upon binding changes

rabbitmq-top

Top like view of RabbitMQ processes in the management UI

sharding

Scale out RabbitMQ with automatic queue sharding

(and many more)

Message Routing

Built-In Exchange Types

- ❖ **Direct**

String matching on the routing key

- ❖ **Fanout**

No routing key, messages delivered to all bound queues

- ❖ **Topic**

Pattern matching in the routing key

- ❖ **Headers**

No routing key, value matching in the headers property

Topic Exchange Binding Keys

Routing Key: **namespace.delimited.keys**

#

Receive all messages

namespace.#

Receive all messages in namespace

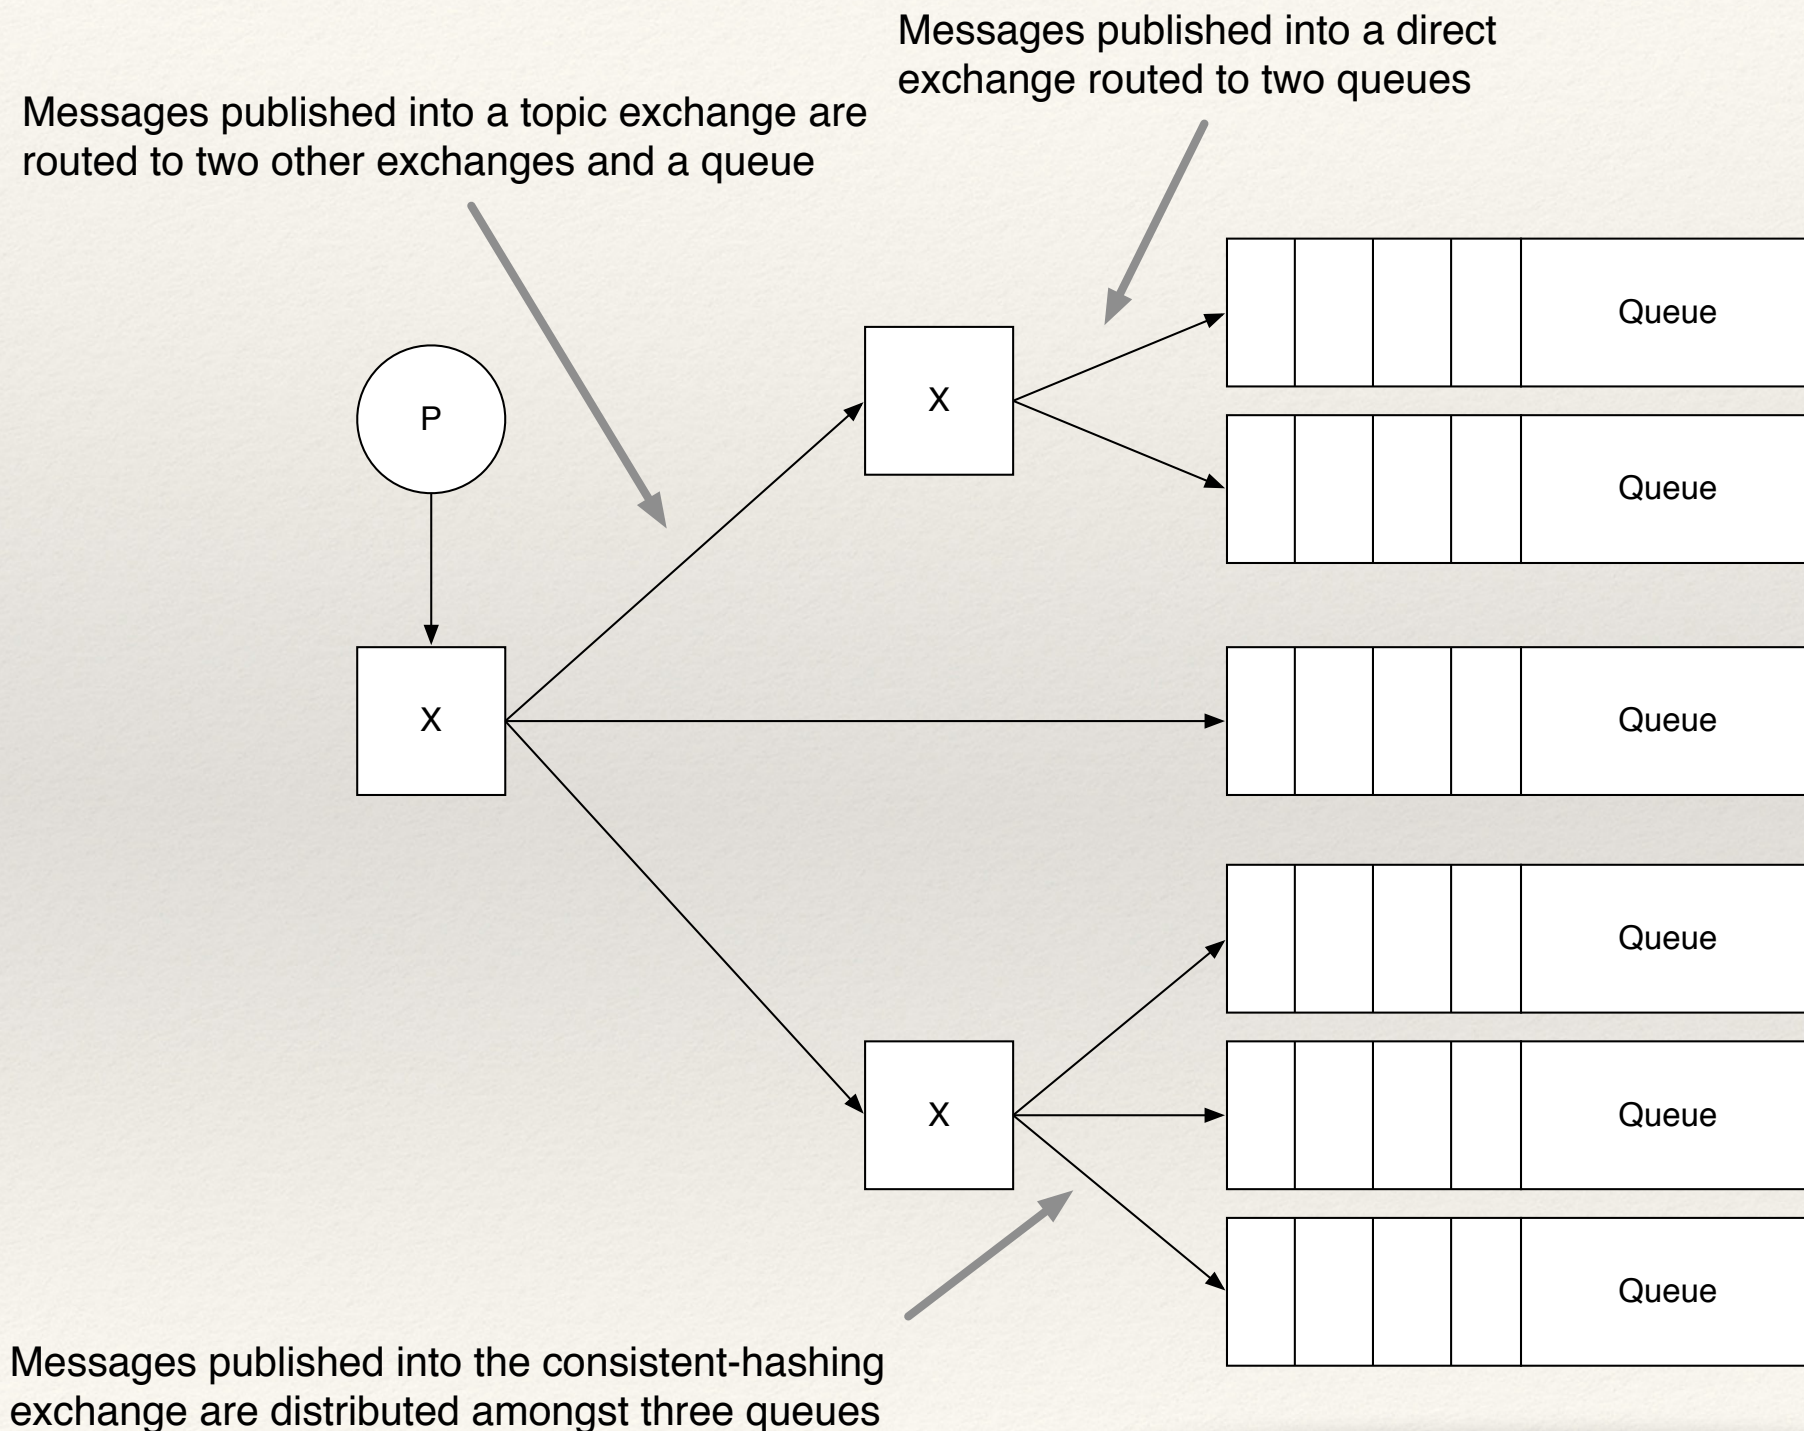
namespace.delimited.*

Receive all namespace.delimited messages

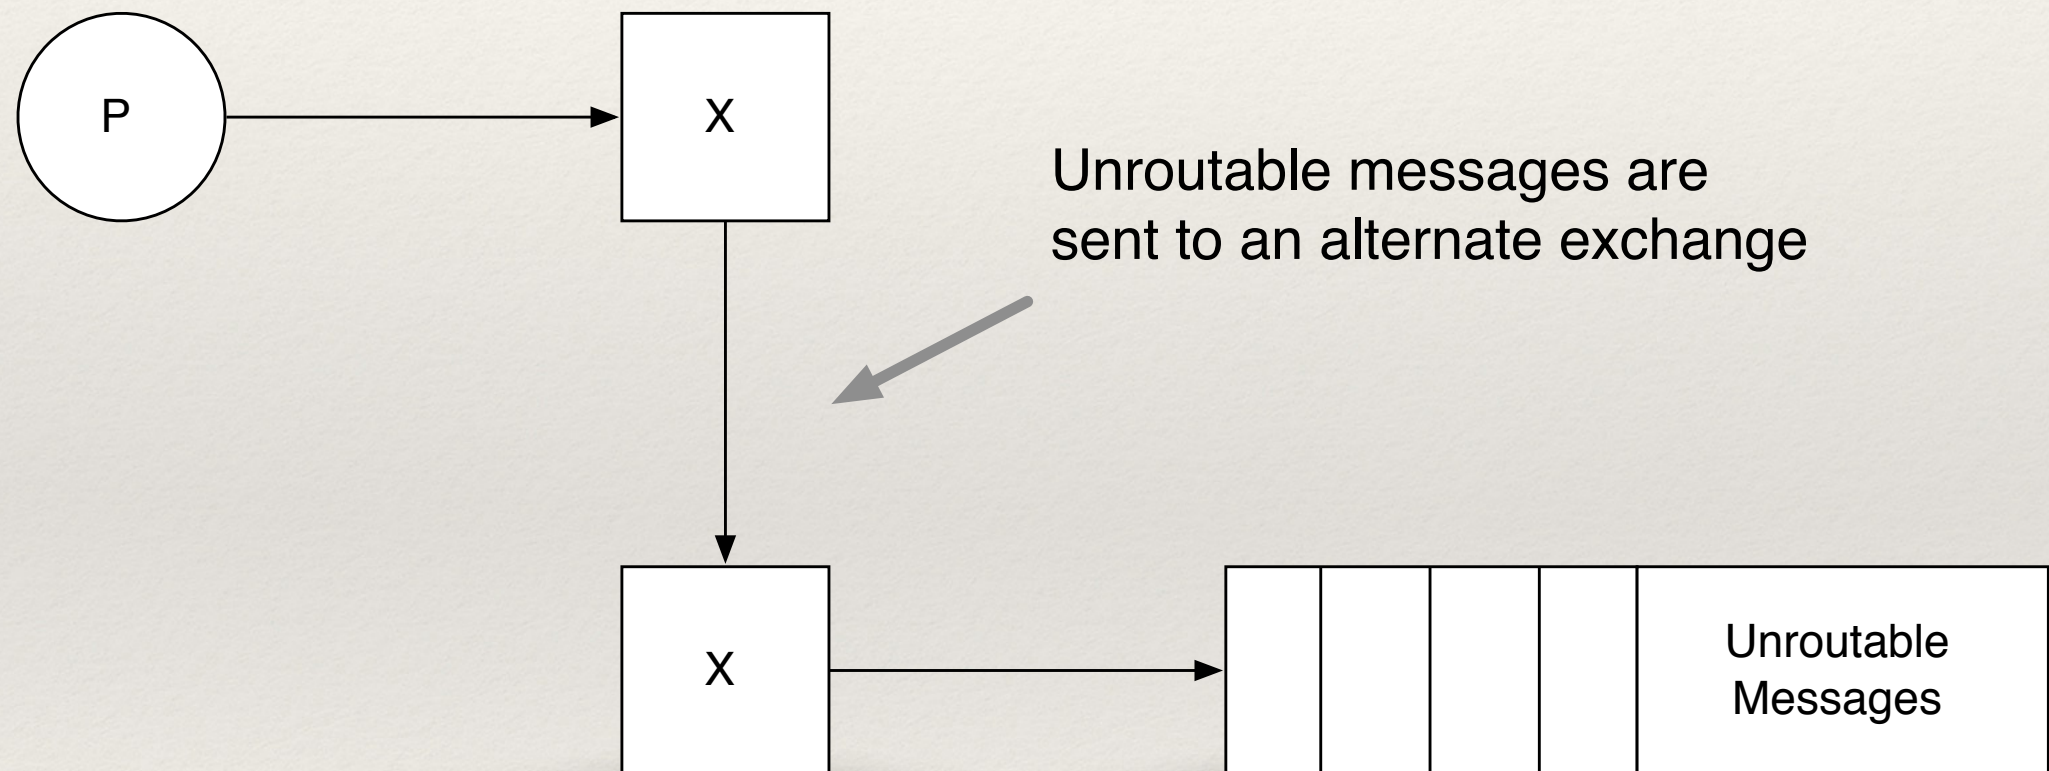
namespace.*.keys

Receive all namespace messages ending with keys

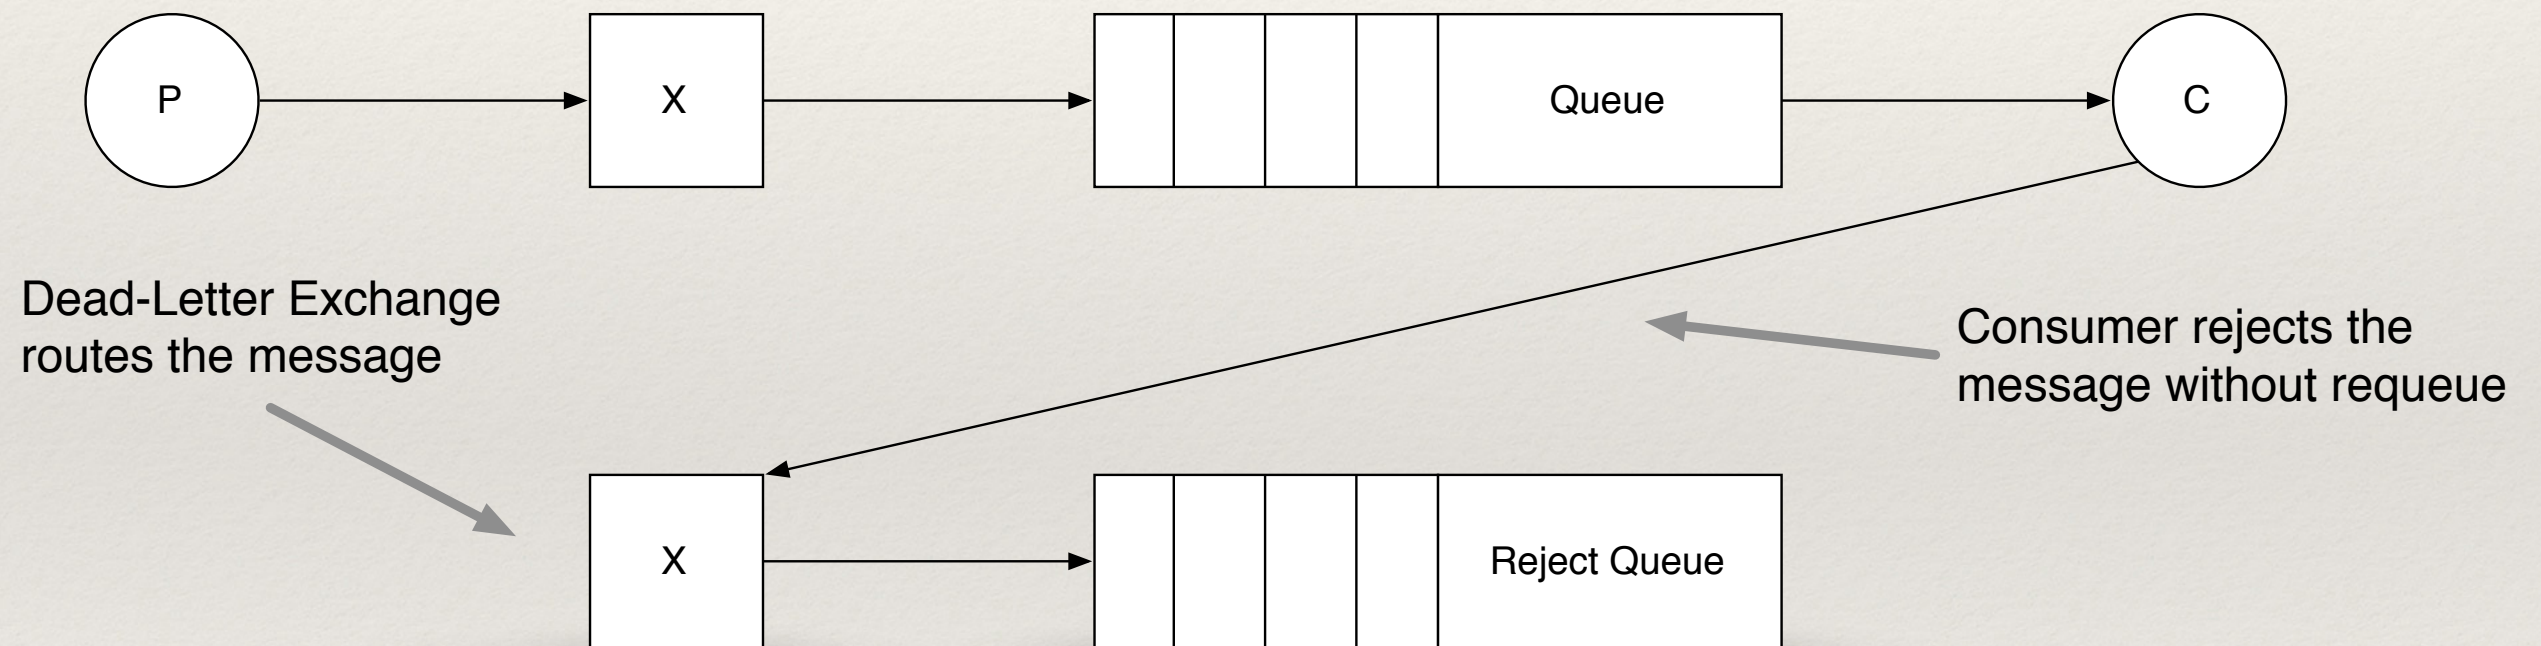
Exchange to Exchange Binding



Alternate Exchanges



Dead Letter Exchanges



Exchange Plugins

Consistent Hashing

Distribute messages via hashed value of routing key

Event

Publishes messages on AMQP events such as queue creation

Random

Distribute messages across all bound queues randomly

PostgreSQL LISTEN

Subscribes to and publishes PostgreSQL notifications

Recent History

Sends the last n messages to any newly bound queue

Reverse Topic

Allows for routing patterns at publish time and not via binding

Riak Storage

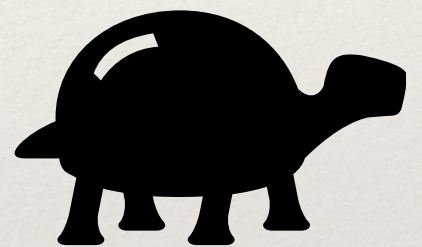

Stores messages published through the exchange into Riak

Script Exchange

Calls out to external scripts for message routing

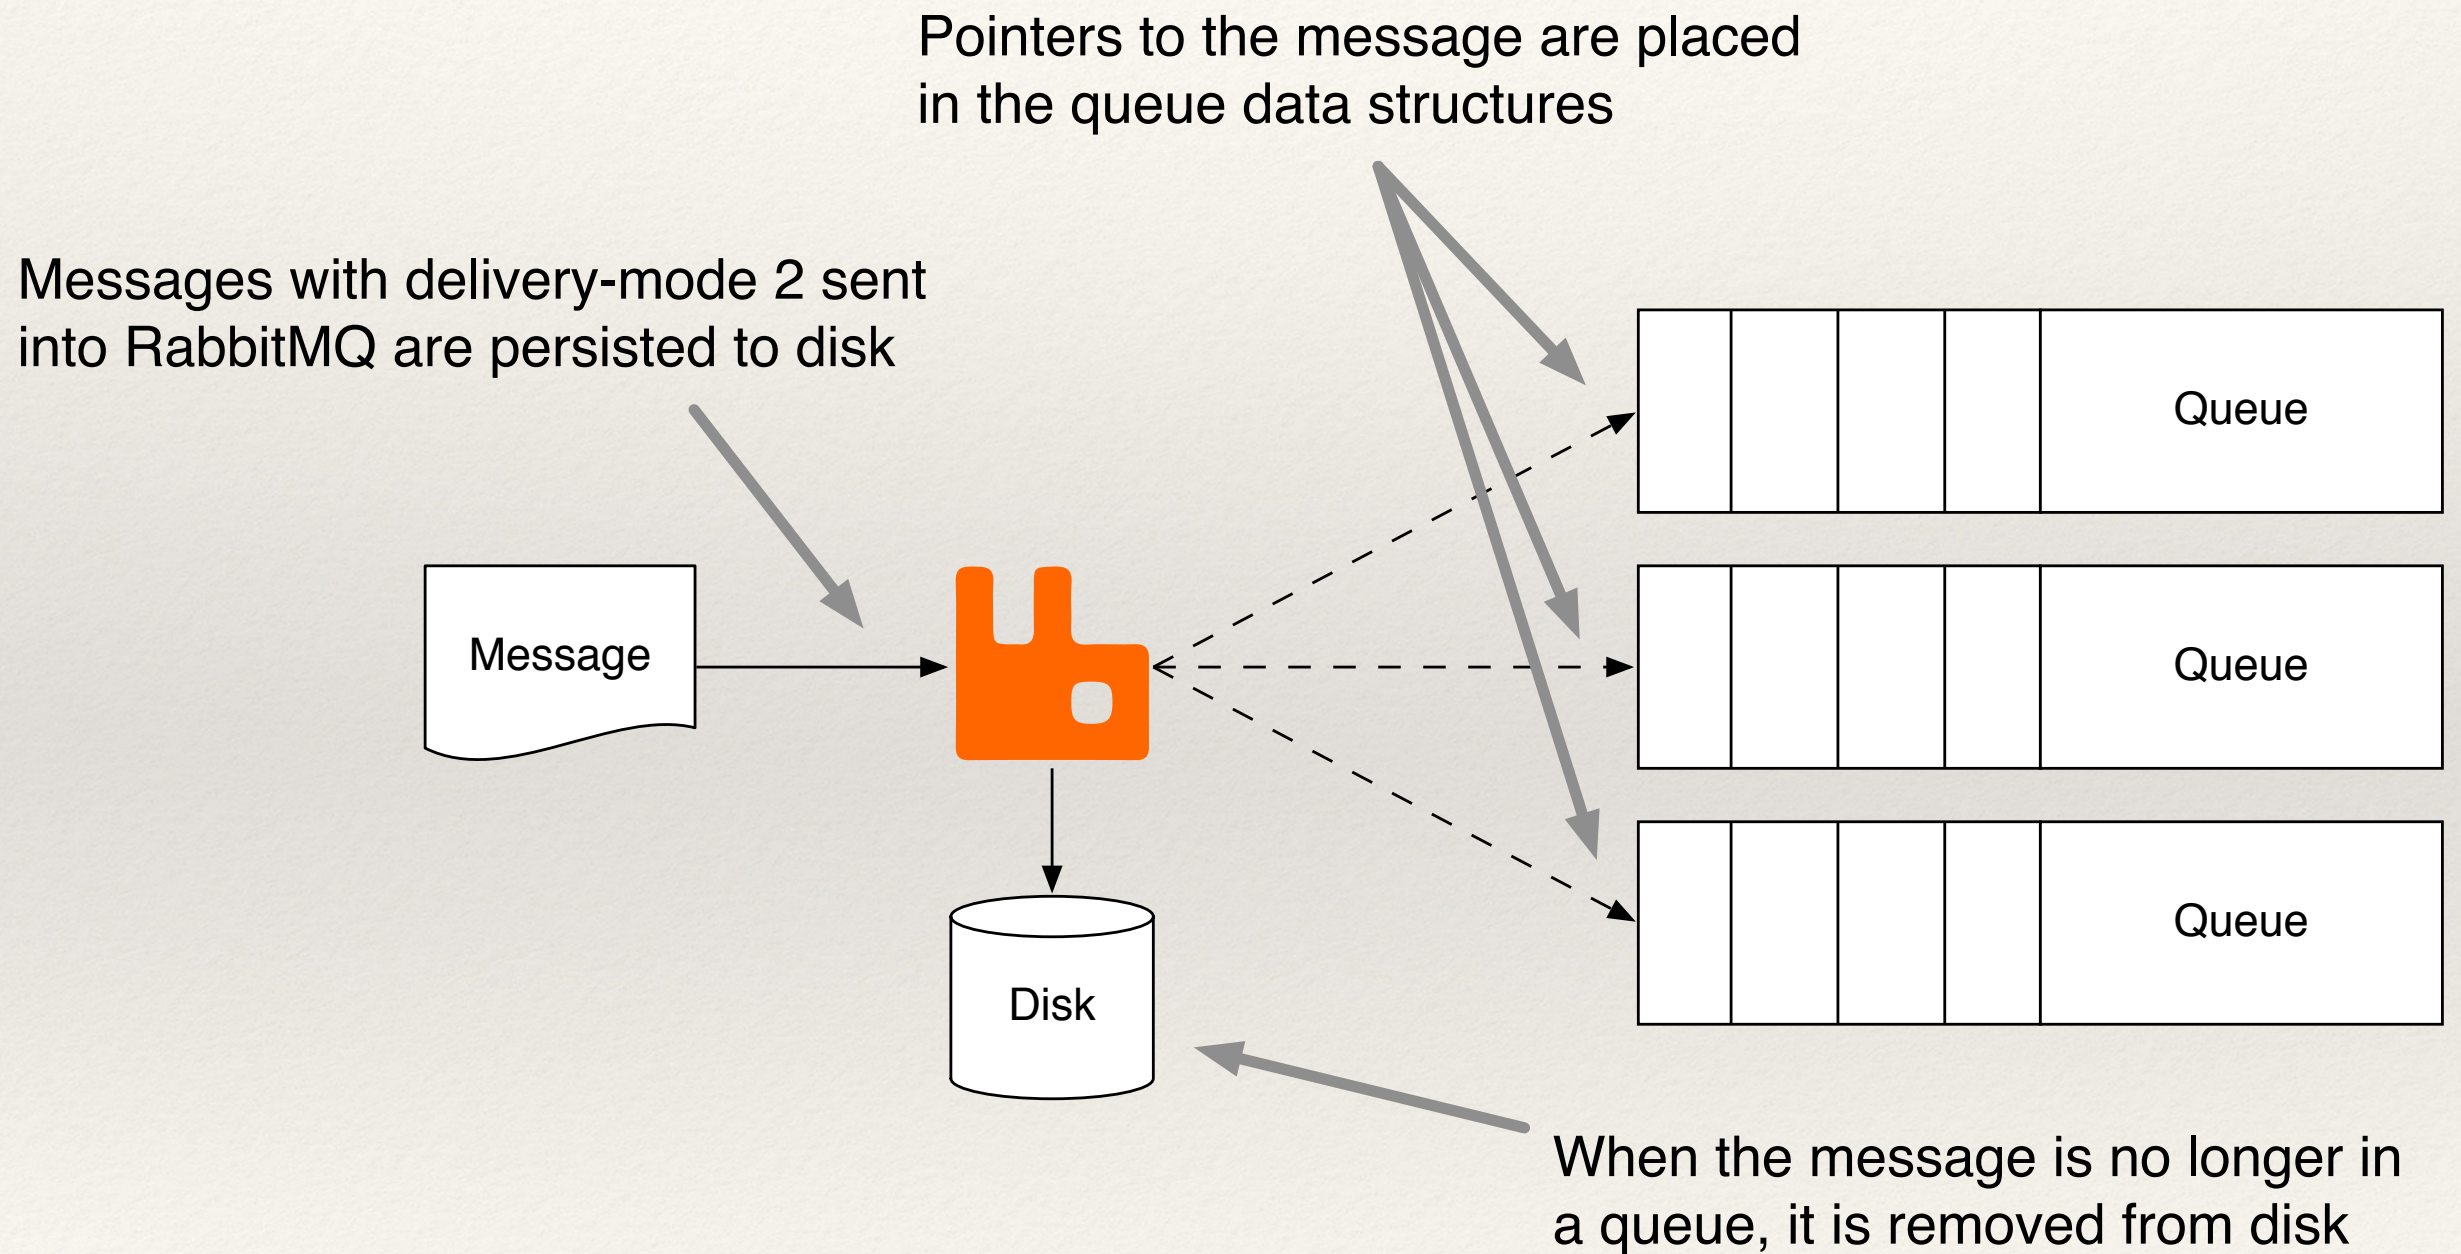
Performance Considerations

Publishing Performance Scale

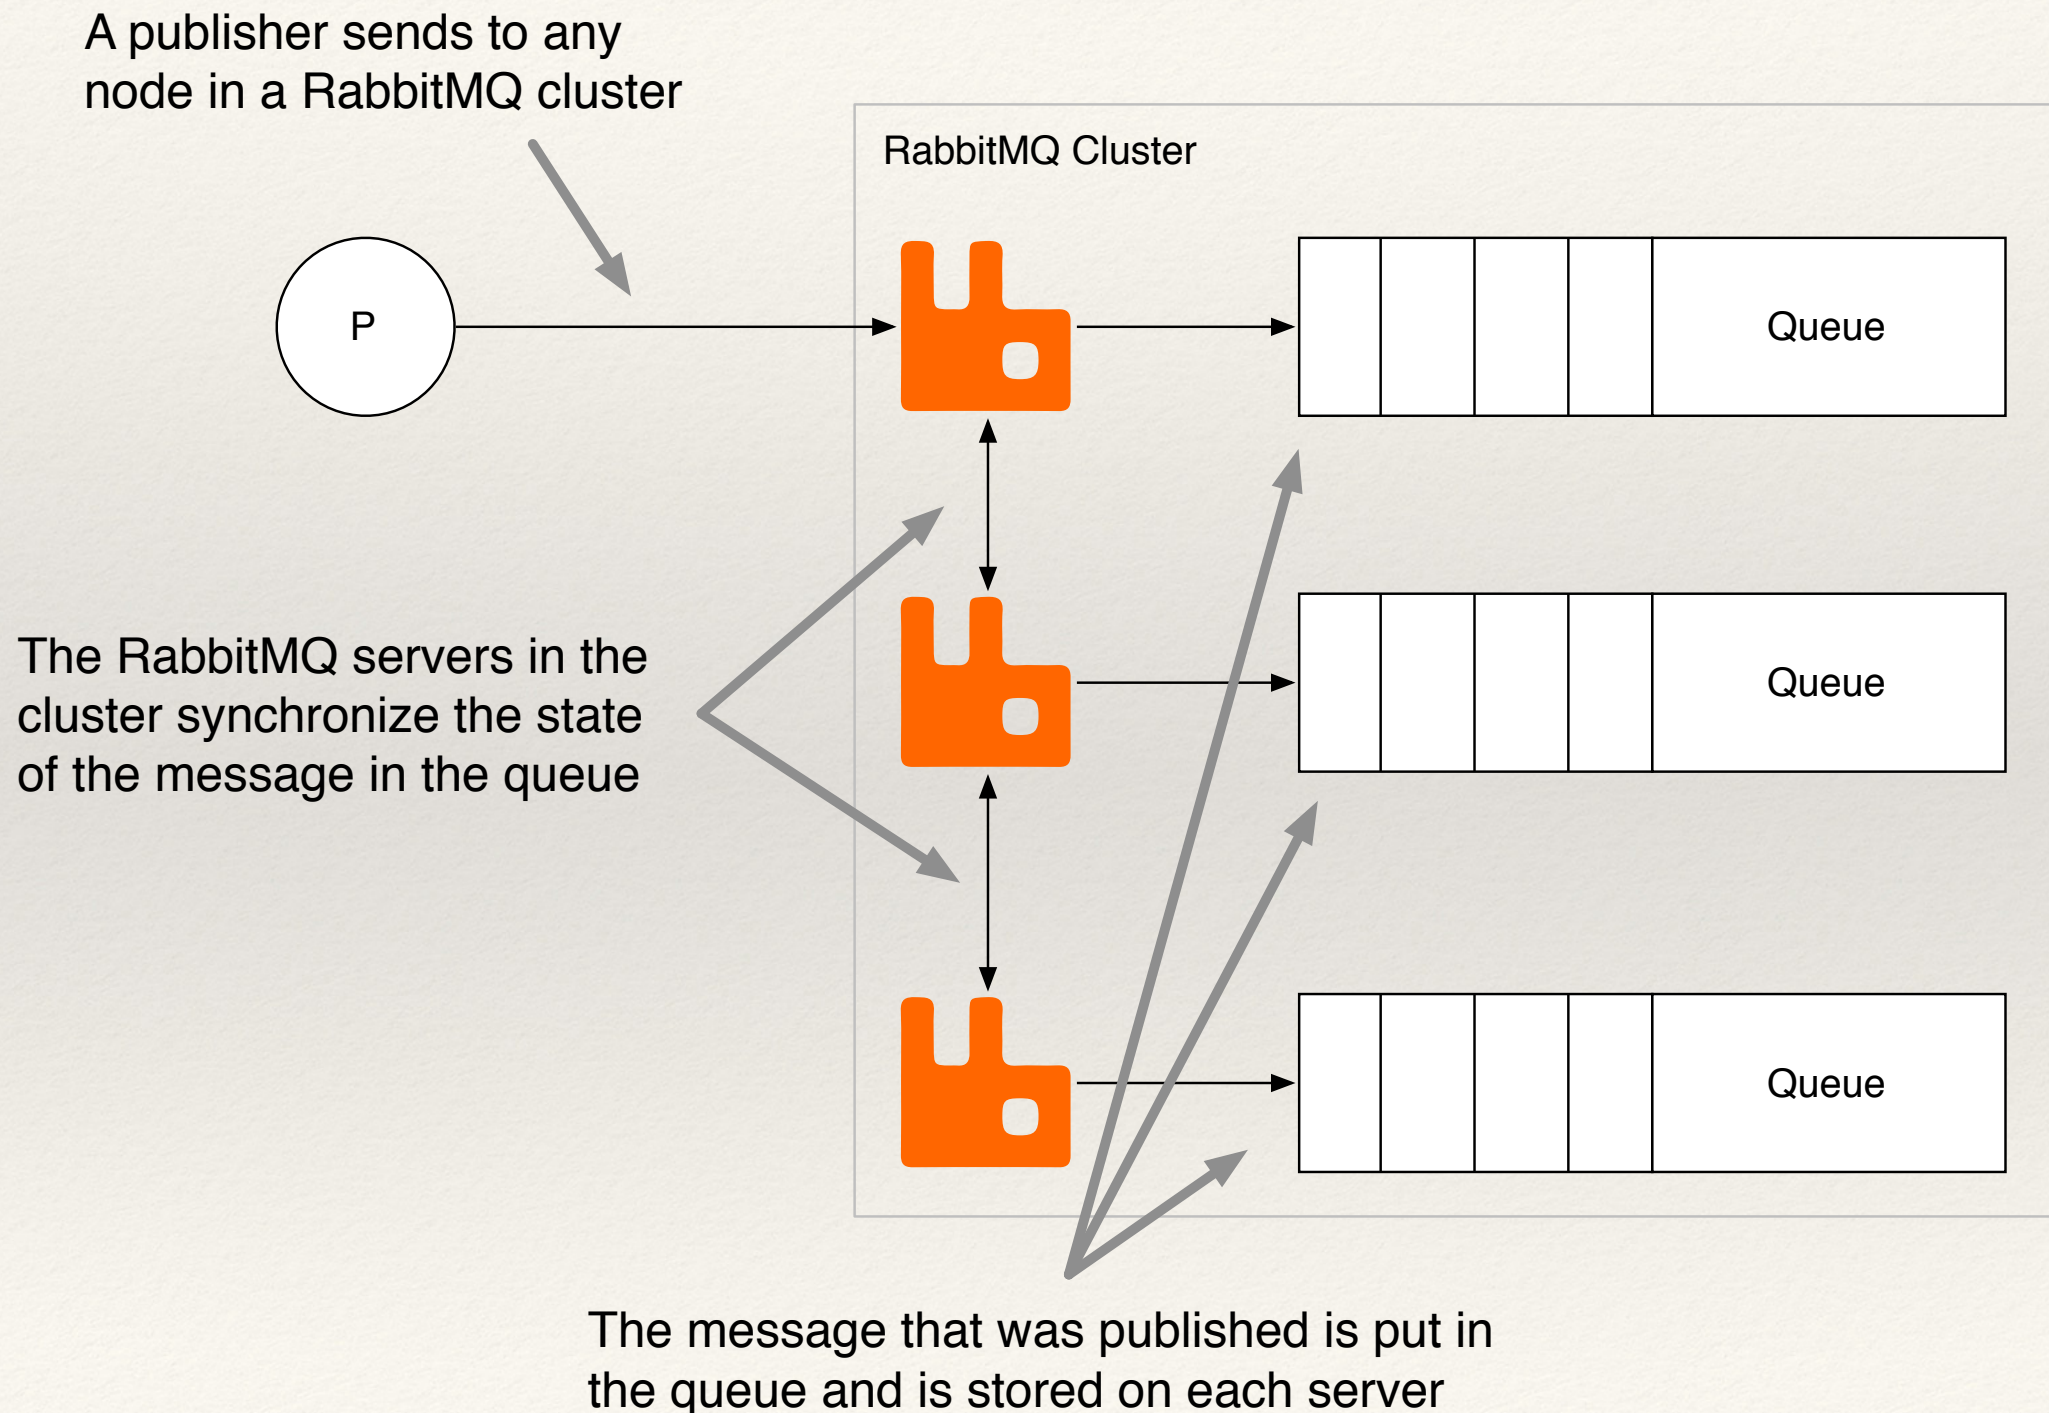


No guarantees
Notification on failure
Publisher confirms
Alternate exchanges
HA Queues
Transactions
HA Queues w/ Transactions
Persisted Messages


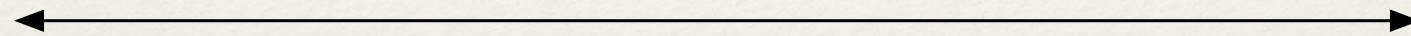
Persisted Messages IO Model



HA Queues & Performance



Consumer Performance Scale



Getting Messages
Consuming and using Transactions
Consuming with Acknowledgements
Consume with "No Ack Mode" enabled
Consuming with Acknowledgements and QoS > 1

Operational Concerns

Trending & Monitoring

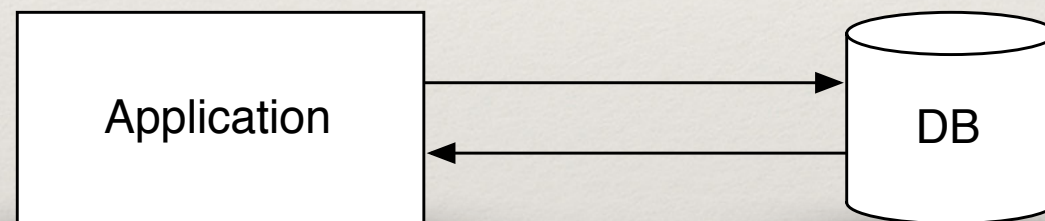
- ❖ RabbitMQ Management Plugin provides internal stats
 - ❖ *Queue depths, connection counts, throughput, memory usage, etc*
- ❖ Monitor with common tools such as Nagios or Sensu to services such as Boundary and NewRelic
- ❖ Stream based monitoring with Riemann for anomaly detection

Thoughts on Configuration

- ❖ Use configuration management!
 - ❖ Even for exchanges, queues, and bindings if possible
 - ❖ Helpful for disaster recovery
- ❖ Use RabbitMQ's Policies when possible
 - ❖ Exchanges and queues are immutable
 - ❖ Deleting and redeclaring for changes can be disruptive

Sample Usage Patterns

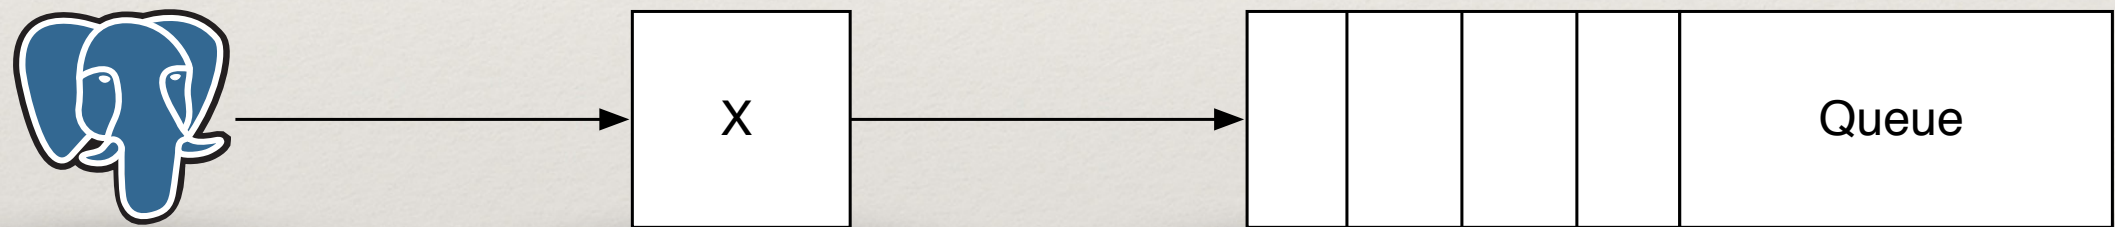
Evolve Tightly-Coupled Applications



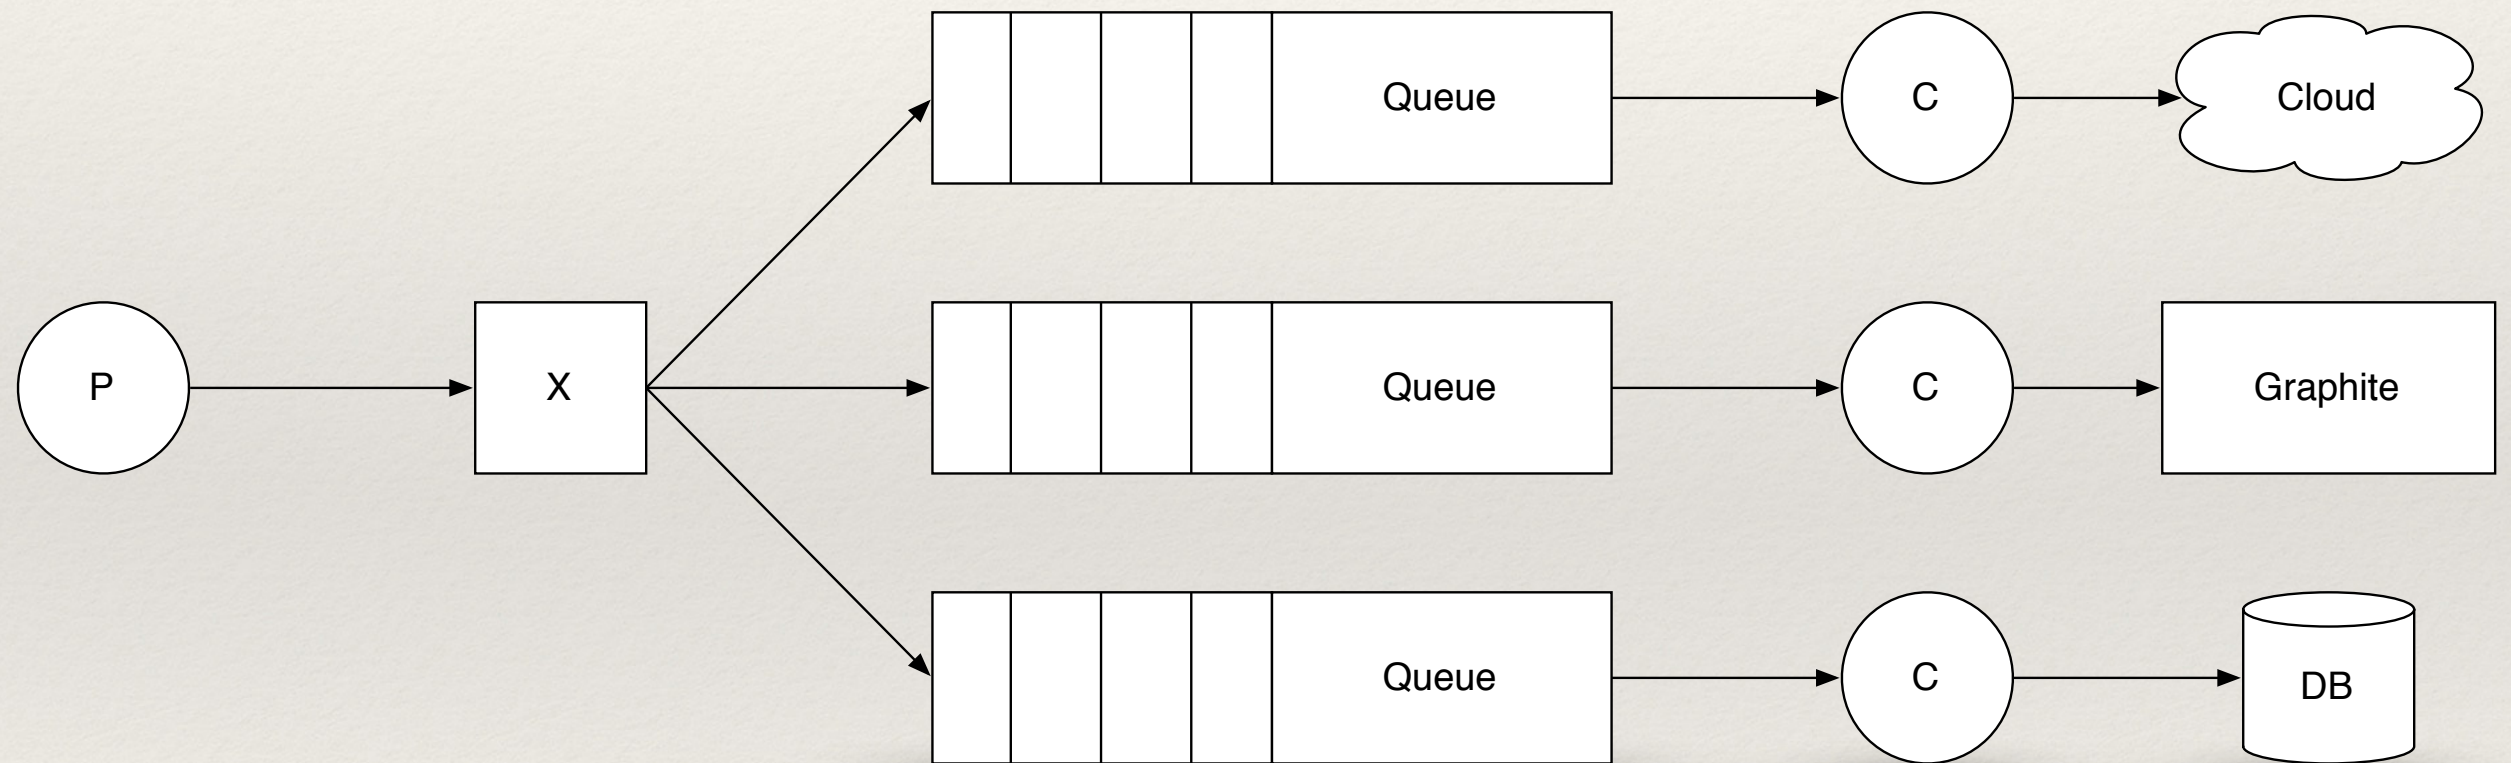
Decoupling Database Writes



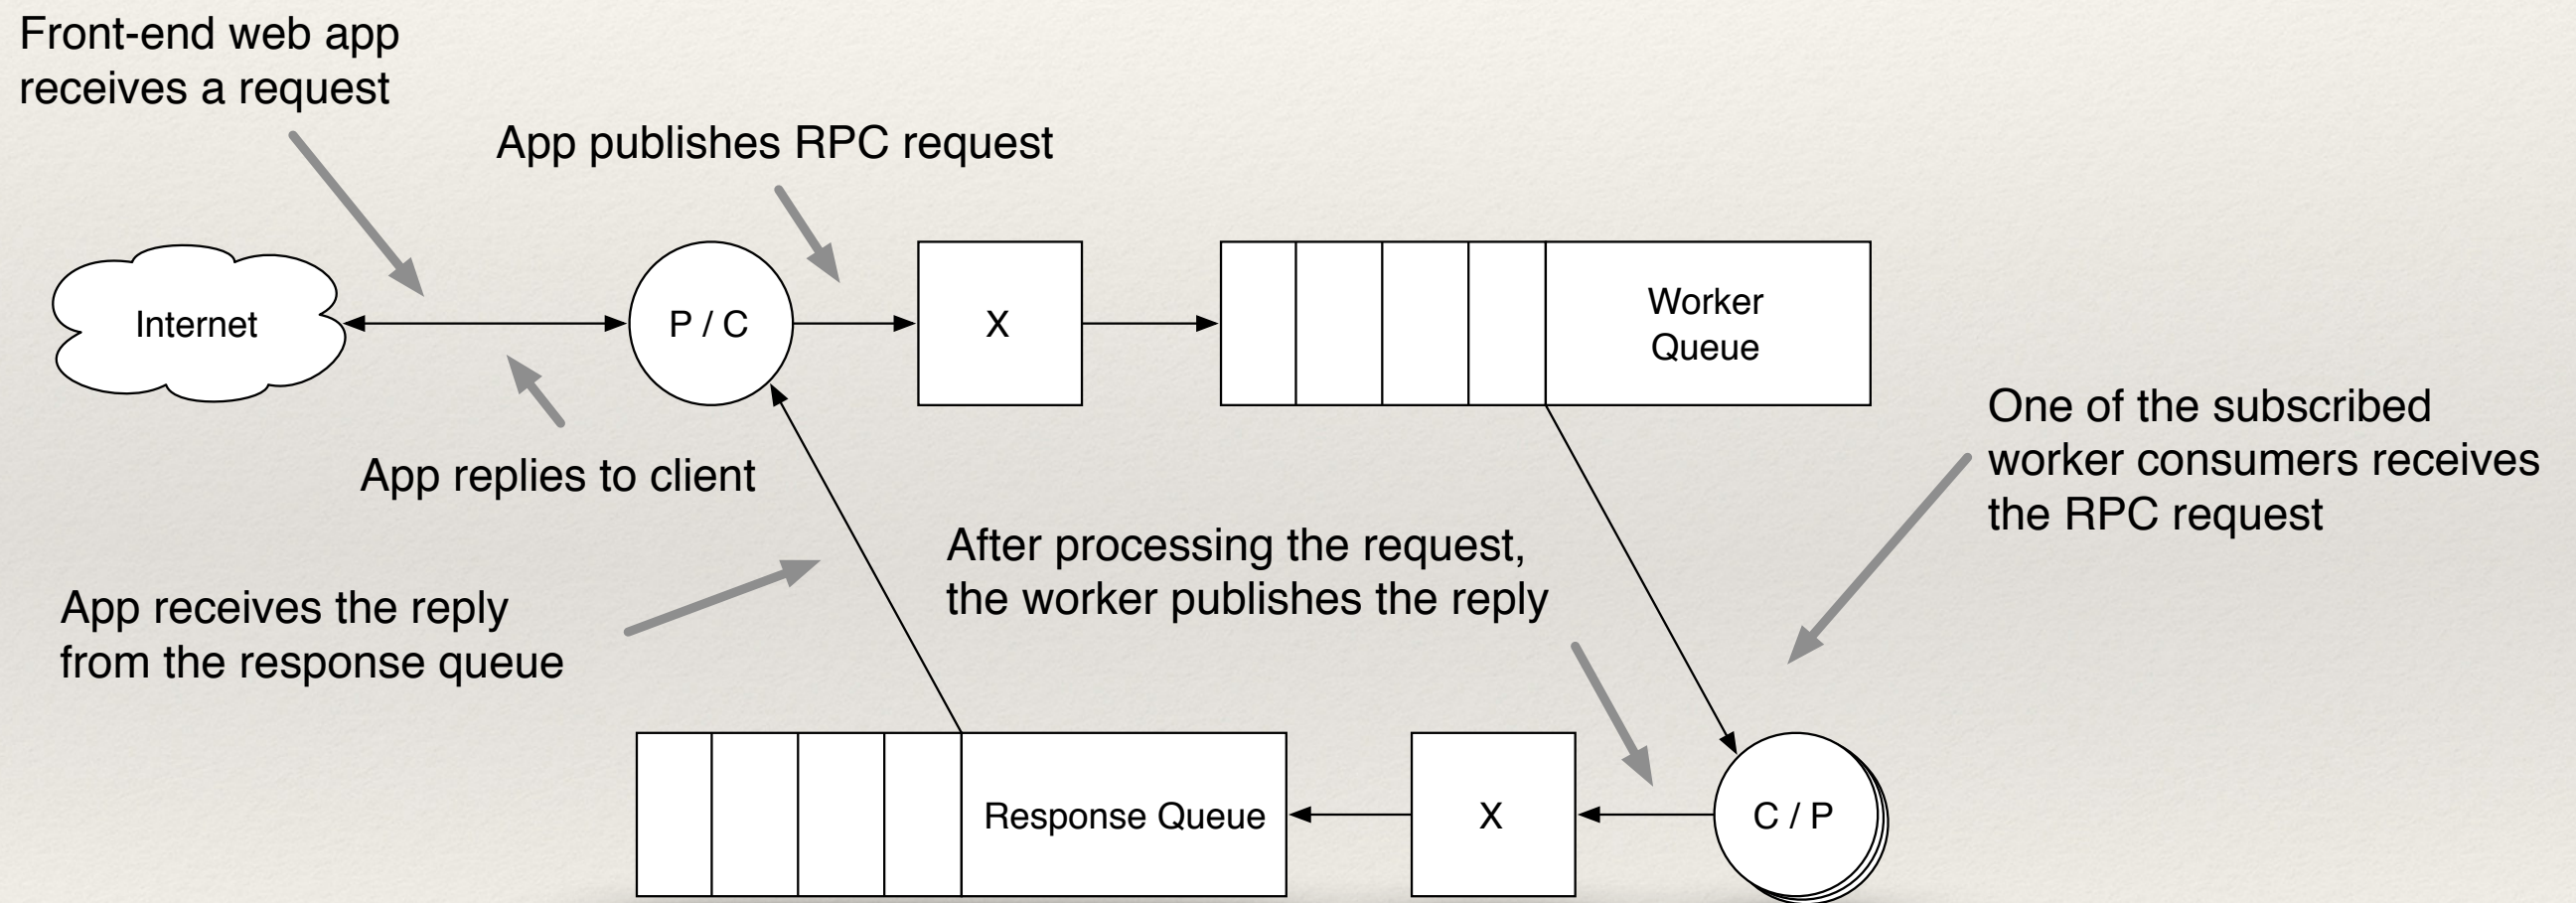
Listen for Database Notifications



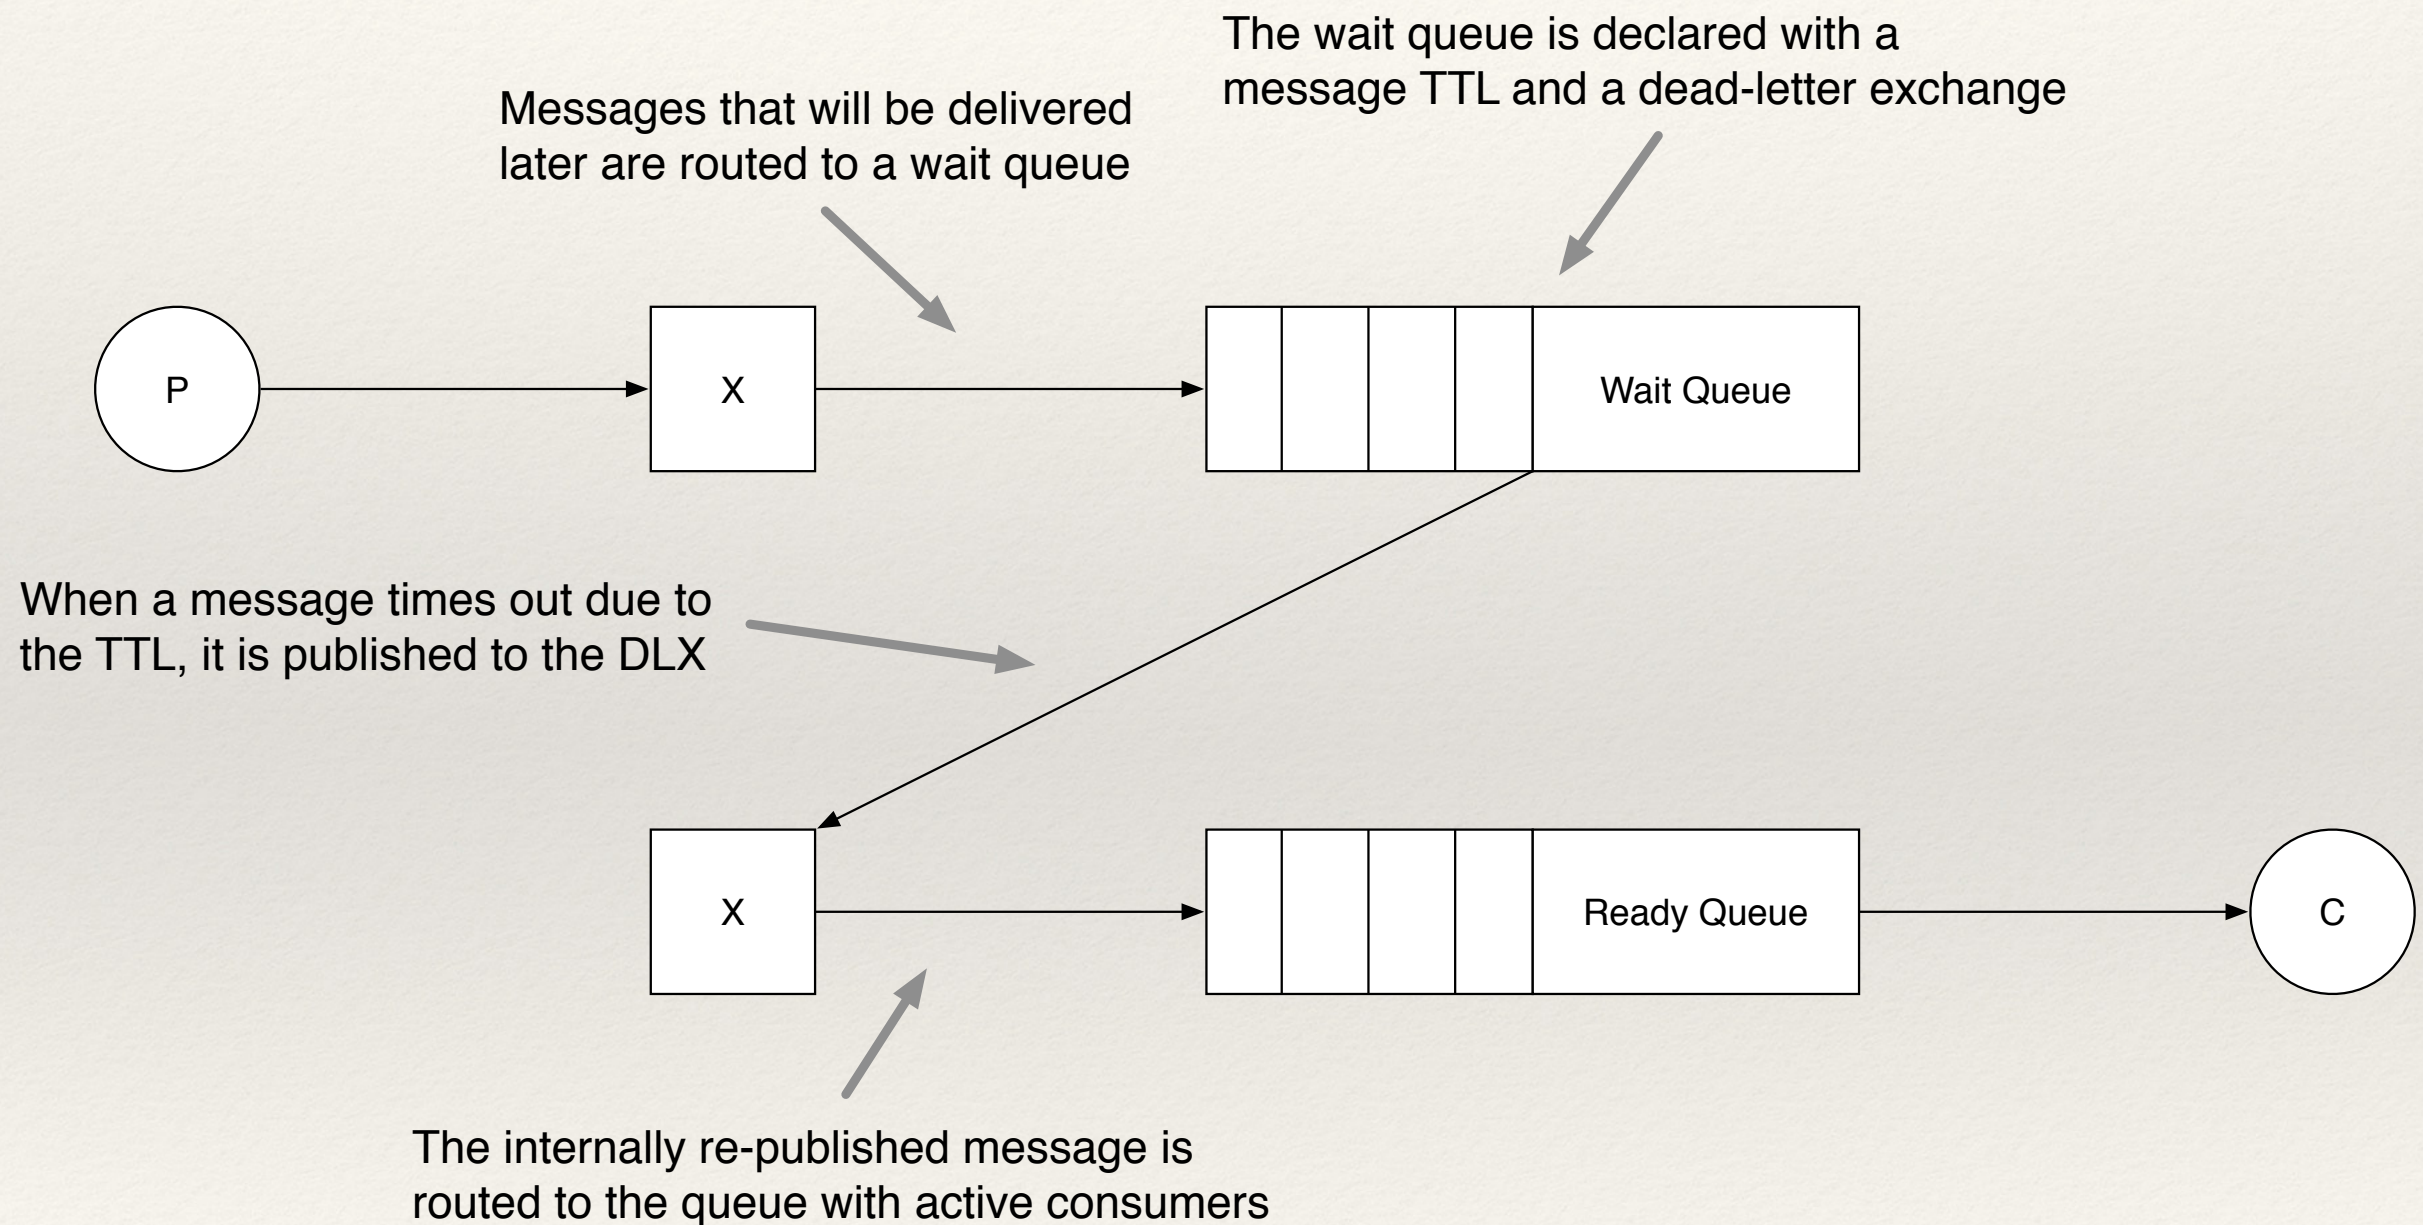
Multi-Purposed Messages



RPC

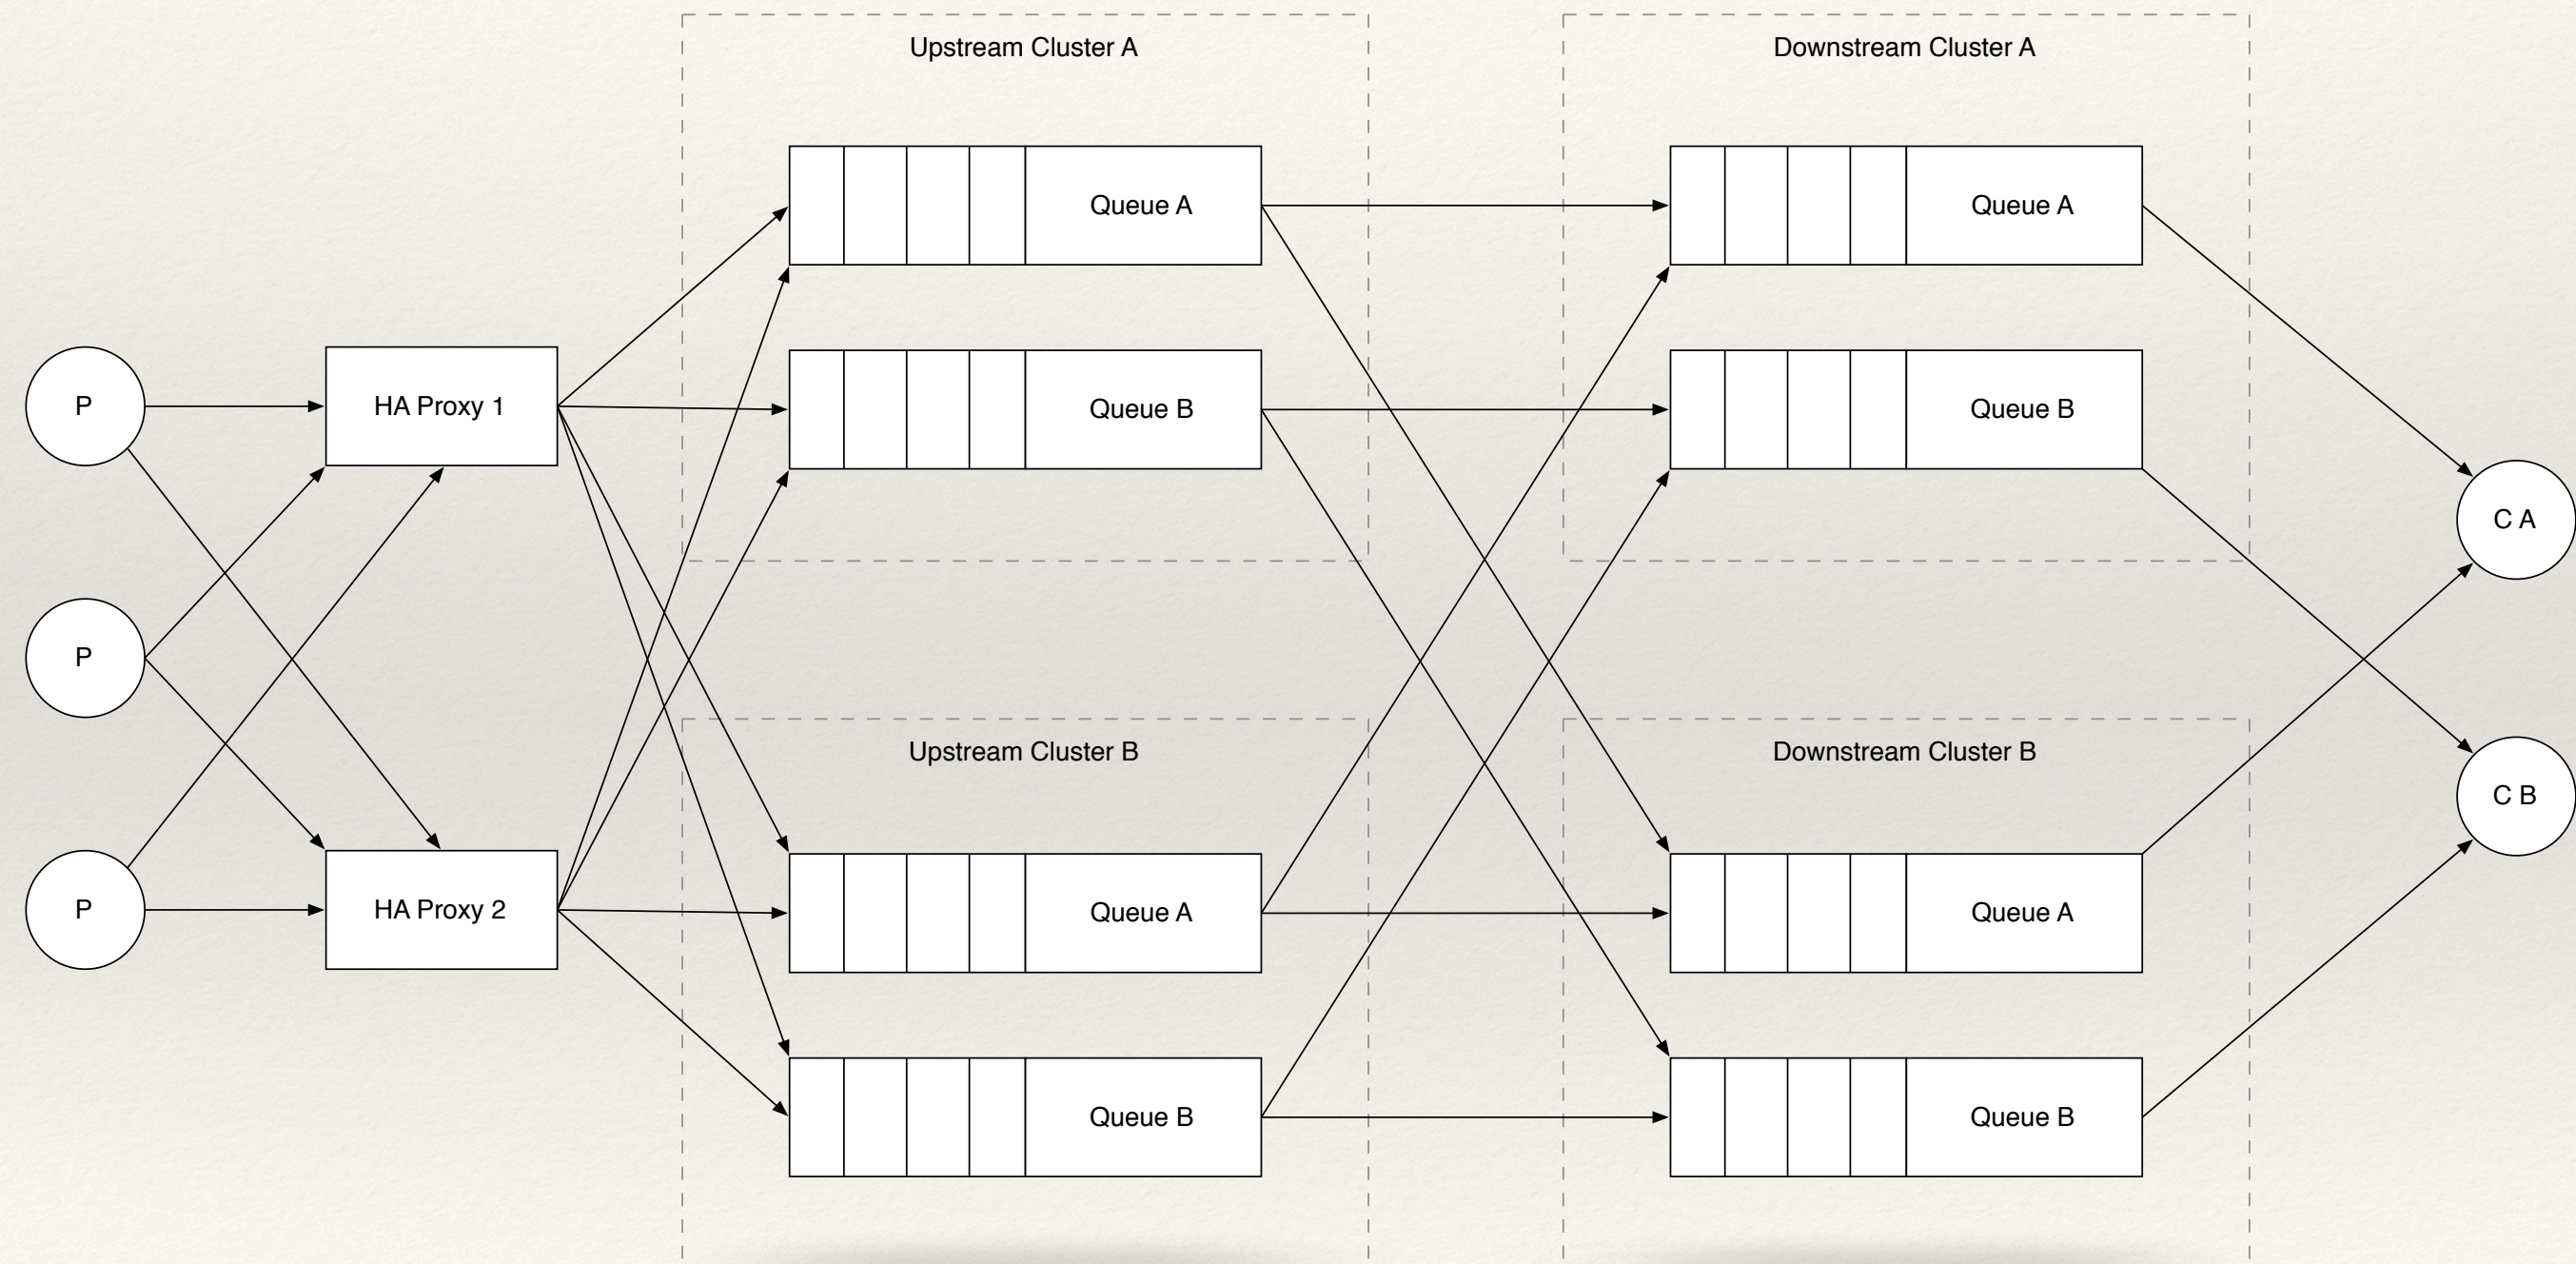


Delayed Messages



WARNING: This pattern is a hack and can cause undesirable situations should the wait queue overwhelm the system!

HA RabbitMQ with Federation



Questions?